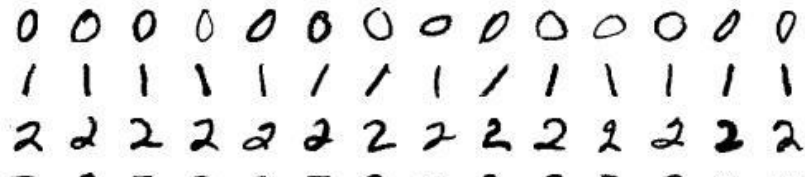



# PyTorch

# Libraries to install

- Numpy
- Torch
- Sklearn
- How?
  - `pip install numpy`, `pip install scikit-learn`, `pip install torch`
  - `conda install numpy`, `conda install scikit-learn`, `conda install scikit-learn`
- If pip is not installed:
  - `python -m pip install --upgrade pip`

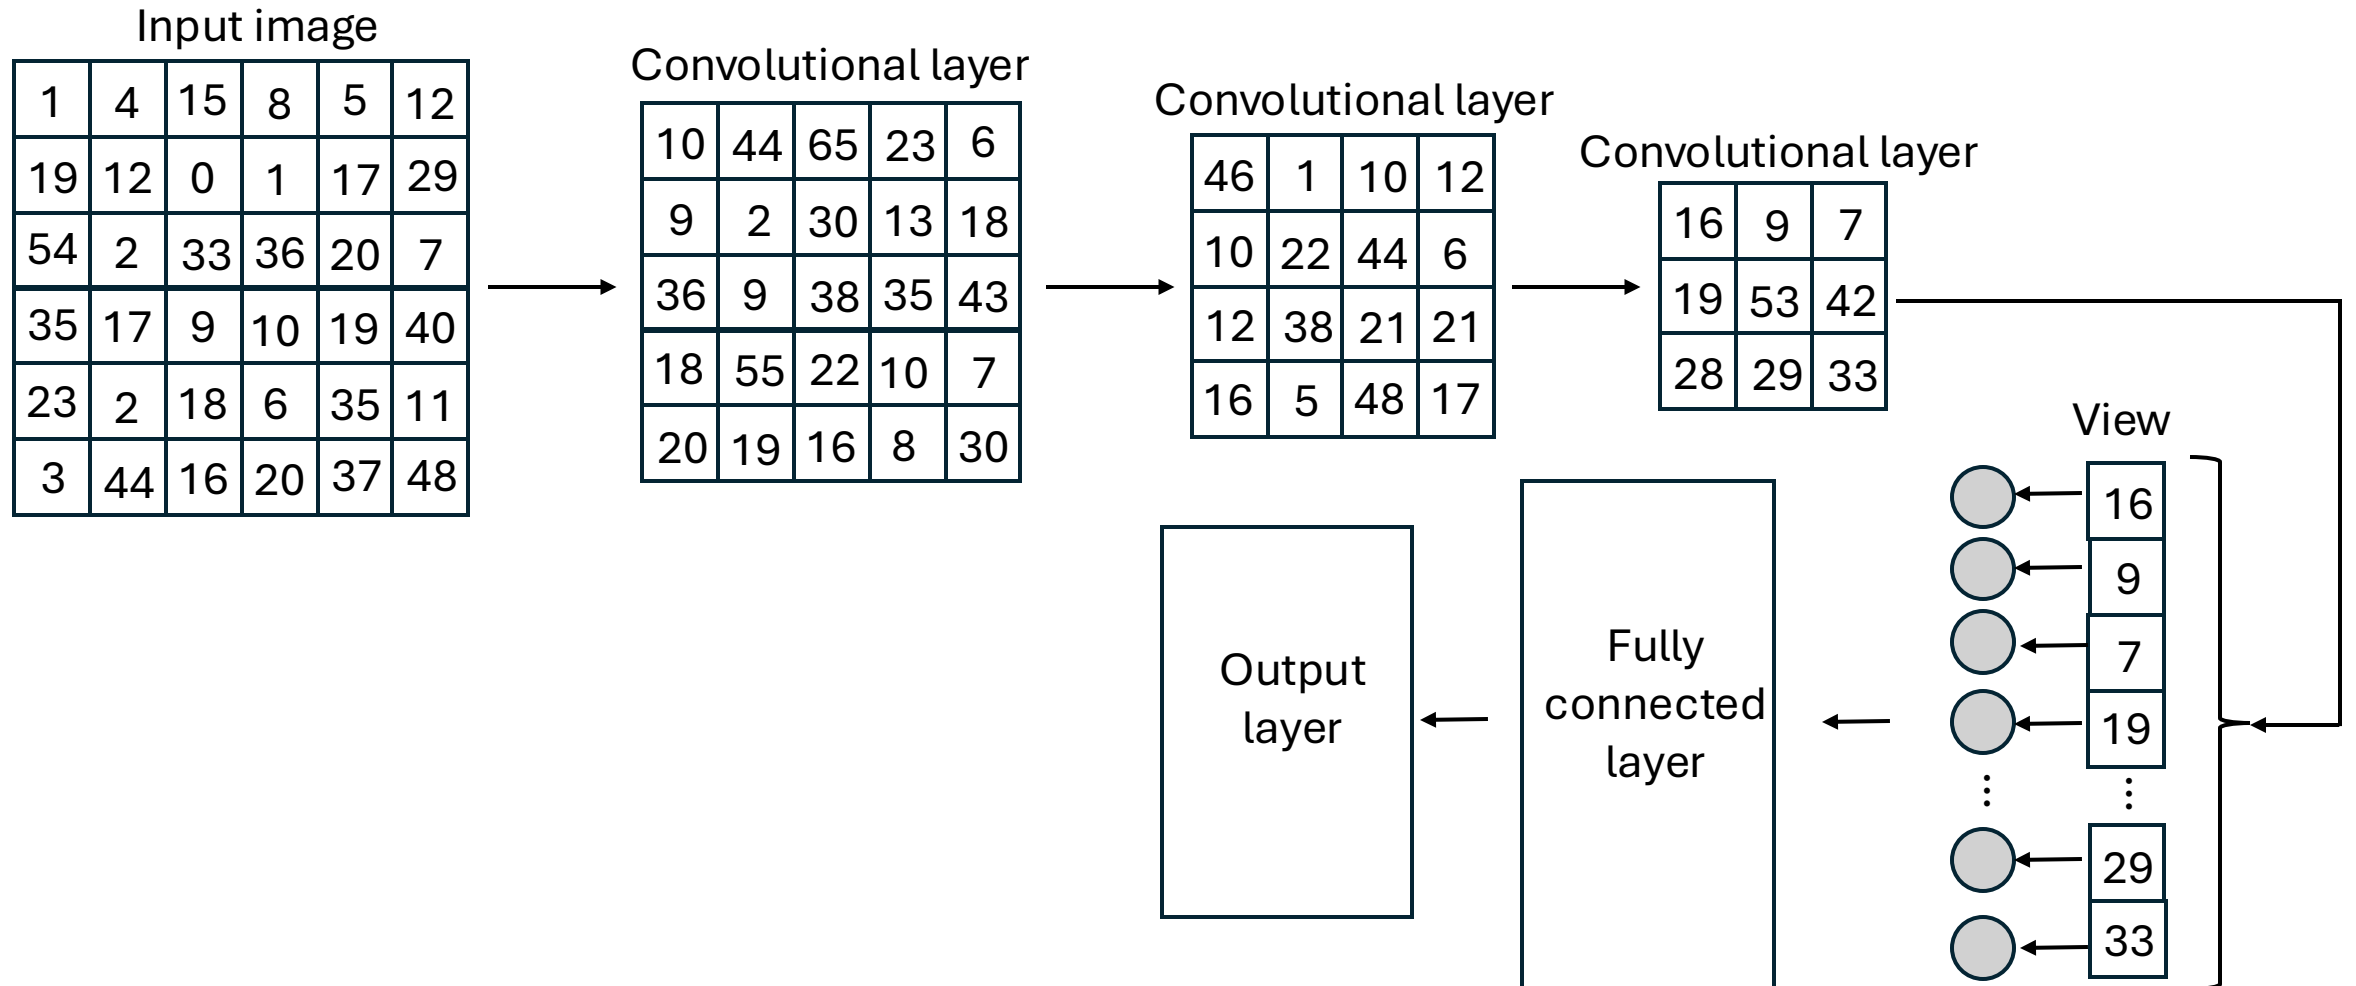
# Problem

- MNIST Database: Handwritten Digits.
  - The goal is to build a neural network that can **recognize digits from 0 to 9**.
  - We will use **two types of networks**:
    - CNN (Convolutional Neural Network).
    - Fully-connected Neural Network.
- 
- A 4x15 grid of handwritten digits from the MNIST database. The digits are arranged in four rows and fifteen columns. The first row contains mostly 0s and 1s. The second row contains mostly 1s and 2s. The third row contains mostly 2s and 3s. The fourth row contains mostly 3s and 4s. The handwriting is varied, showing different styles and orientations.



# Convolutional layers

- The resulting feature map size decreases with each convolutional layer.



# Max Pooling

MaxPooling



2x2




46	1	10	12
10	22	44	6
12	38	21	21
16	5	48	17



46	44
38	48

# Dropout

- Assigns a weight of 0 (no contribution) to the specified percentage of neurons.
- Example:



# Hands-on exercises

- Instead of building a model to classify all 10 digits, convert the task into a **binary classification** problem — for example, distinguishing between **0 and 8**. To do this:
  - **Select only the images** corresponding to digits **0** and **8** from the dataset.
  - **Modify the neural network accordingly** to handle binary output.

# Hands-on exercises

- To use BCELoss() instead of CrossEntropyLoss()
  - `criterion = nn.BCELoss()`
    - `onehot_encoder = OneHotEncoder(sparse_output=False)`
    - `labels_oh = onehot_encoder.fit_transform(labels.reshape(-1,1))`
    - `labels = labels_oh.astype('float32')`
  - `self.sig = nn.Sigmoid()`
  - `bal_acc =`  
`balanced_accuracy_score(onehot_encoder.inverse_transform(np.reshape(labels_test,[labels_test.shape[0],labels_test.shape[2]])),predictions)`